

```
In [ ]: # Deep Learning
        # Practical No : 01
```

```
In [ ]: import tensorflow as tf
        from sklearn.model_selection import train_test_split
        from sklearn.preprocessing import StandardScaler
        from sklearn import metrics
        import pandas as pd
        import numpy as np
        import matplotlib.pyplot as plt
        import seaborn as sns
        %matplotlib inline
        from tqdm.notebook import tqdm
        import warnings
        warnings.filterwarnings("ignore")
```

```
In [ ]: boston = tf.keras.datasets.boston_housing
```

```
In [ ]: dir(boston)
```

```
Out[ ]: ['__builtins__',
        '__cached__',
        '__doc__',
        '__file__',
        '__loader__',
        '__name__',
        '__package__',
        '__path__',
        '__spec__',
        'load_data']
```

```
In [ ]: boston_data = boston.load_data()
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/
boston_housing.npz
57026/57026 [=====] - 0s 0us/step
```

```
In [ ]: (x_train, y_train), (x_test, y_test) = tf.keras.datasets.boston_housing.load_data(p
```

```
In [ ]: x_train.shape, y_train.shape, x_test.shape, y_test.shape
```

```
Out[ ]: ((404, 13), (404,)), (102, 13), (102,))
```

```
In [ ]: scaler = StandardScaler()
```

```
In [ ]: x_train_scaled = scaler.fit_transform(x_train)
        x_test_scaled = scaler.transform(x_test)

        y_train_scaled = scaler.fit_transform(y_train.reshape(-1, 1))
        y_test_scaled = scaler.transform(y_test.reshape(-1, 1))
```

```
In [ ]: model = tf.keras.models.Sequential([
        tf.keras.layers.Input(shape=(13), name='input-layer'),
```

```
tf.keras.layers.Dense(100, name='hidden-layer-2'),
tf.keras.layers.BatchNormalization(name='hidden-layer-3'),
tf.keras.layers.Dense(50, name='hidden-layer-4'),
tf.keras.layers.Dense(1, name='output-layer')
])
```

```
In [ ]: model.summary()
```

Model: "sequential\_10"

Layer (type)	Output Shape	Param #
=====		
hidden-layer-2 (Dense)	(None, 100)	1400
hidden-layer-3 (BatchNormalization)	(None, 100)	400
hidden-layer-4 (Dense)	(None, 50)	5050
output-layer (Dense)	(None, 1)	51
=====		
Total params: 6901 (26.96 KB)		
Trainable params: 6701 (26.18 KB)		
Non-trainable params: 200 (800.00 Byte)		

```
In [ ]: model.compile(
    optimizer='adam',
    loss='mse',
    metrics=['mae']
)
```

```
In [ ]: history = model.fit(x_train, y_train, batch_size=32, epochs=20, validation_data=(x_
```

Epoch 1/20  
13/13 [=====] - 1s 18ms/step - loss: 574.8789 - mae: 22.3634 - val\_loss: 2224.5266 - val\_mae: 44.4160

Epoch 2/20  
13/13 [=====] - 0s 7ms/step - loss: 550.9838 - mae: 22.0251 - val\_loss: 705.3503 - val\_mae: 24.4563

Epoch 3/20  
13/13 [=====] - 0s 7ms/step - loss: 530.6443 - mae: 21.6095 - val\_loss: 1005.8787 - val\_mae: 30.3659

Epoch 4/20  
13/13 [=====] - 0s 6ms/step - loss: 501.7632 - mae: 21.0424 - val\_loss: 629.7626 - val\_mae: 23.7866

Epoch 5/20  
13/13 [=====] - 0s 7ms/step - loss: 462.1404 - mae: 20.1391 - val\_loss: 658.0805 - val\_mae: 24.4096

Epoch 6/20  
13/13 [=====] - 0s 6ms/step - loss: 399.1111 - mae: 18.7222 - val\_loss: 542.7188 - val\_mae: 22.2311

Epoch 7/20  
13/13 [=====] - 0s 7ms/step - loss: 312.8028 - mae: 16.4244 - val\_loss: 313.3749 - val\_mae: 16.6600

Epoch 8/20  
13/13 [=====] - 0s 6ms/step - loss: 209.3056 - mae: 13.1072 - val\_loss: 347.4641 - val\_mae: 17.6651

Epoch 9/20  
13/13 [=====] - 0s 7ms/step - loss: 112.0523 - mae: 8.9190 - val\_loss: 192.2844 - val\_mae: 12.5691

Epoch 10/20  
13/13 [=====] - 0s 7ms/step - loss: 55.1853 - mae: 5.6105 - val\_loss: 115.9377 - val\_mae: 9.1478

Epoch 11/20  
13/13 [=====] - 0s 6ms/step - loss: 33.7891 - mae: 4.2850 - val\_loss: 29.2040 - val\_mae: 3.8567

Epoch 12/20  
13/13 [=====] - 0s 7ms/step - loss: 36.9304 - mae: 4.4974 - val\_loss: 99.4068 - val\_mae: 8.0226

Epoch 13/20  
13/13 [=====] - 0s 6ms/step - loss: 32.5308 - mae: 4.2865 - val\_loss: 35.6296 - val\_mae: 4.8049

Epoch 14/20  
13/13 [=====] - 0s 6ms/step - loss: 29.5778 - mae: 4.0588 - val\_loss: 41.9396 - val\_mae: 4.5167

Epoch 15/20  
13/13 [=====] - 0s 7ms/step - loss: 29.4901 - mae: 4.0193 - val\_loss: 49.0794 - val\_mae: 5.6884

Epoch 16/20  
13/13 [=====] - 0s 6ms/step - loss: 32.3794 - mae: 4.1094 - val\_loss: 26.8985 - val\_mae: 3.5455

Epoch 17/20  
13/13 [=====] - 0s 7ms/step - loss: 31.6614 - mae: 4.0583 - val\_loss: 33.6451 - val\_mae: 3.8305

Epoch 18/20  
13/13 [=====] - 0s 7ms/step - loss: 26.4994 - mae: 3.7155 - val\_loss: 26.4998 - val\_mae: 3.4068

Epoch 19/20  
13/13 [=====] - 0s 7ms/step - loss: 27.2167 - mae: 3.8110

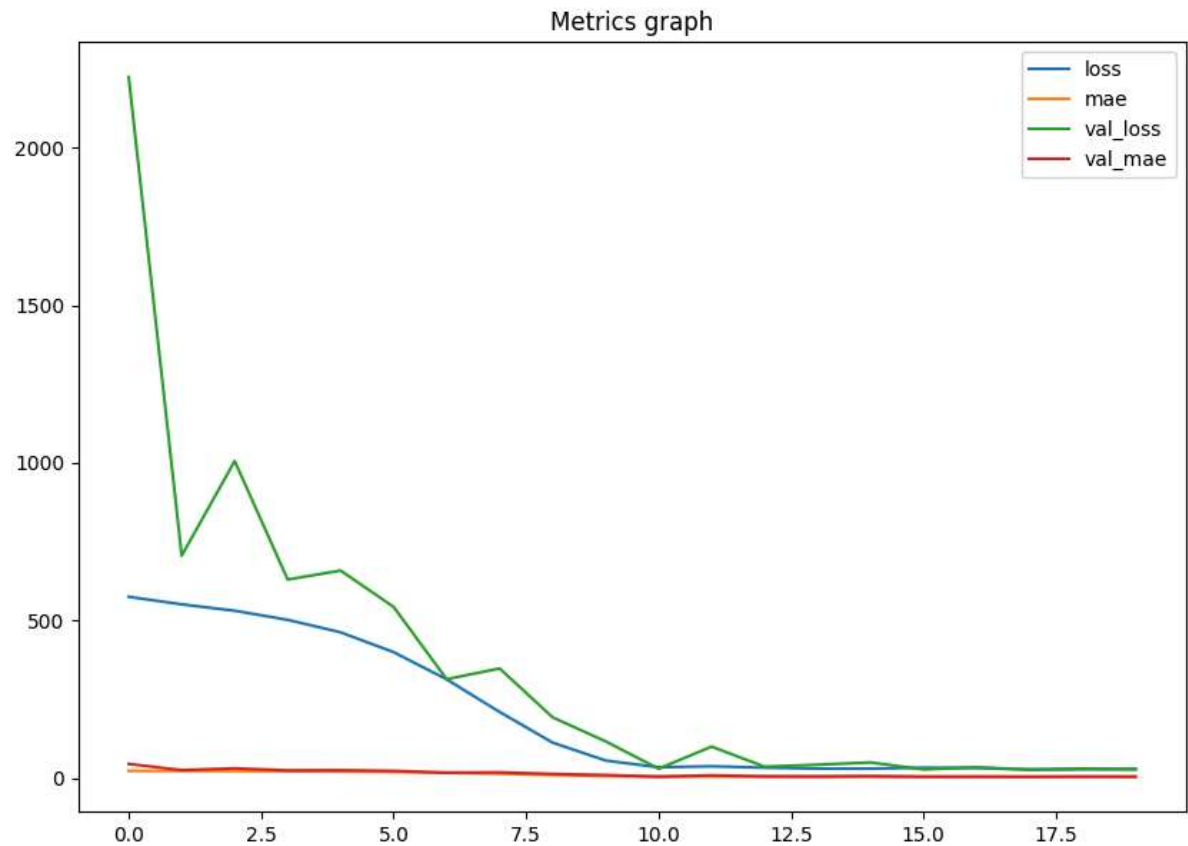
- val\_loss: 29.5473 - val\_mae: 4.1069

Epoch 20/20

13/13 [=====] - 0s 7ms/step - loss: 27.8241 - mae: 3.9048

- val\_loss: 27.6318 - val\_mae: 3.7487


```
In [ ]: pd.DataFrame(history.history).plot(figsize=(10,7))
plt.title("Metrics graph")
plt.show()
```



```
In [ ]: y_pred = model.predict(x_test)
```

4/4 [=====] - 0s 4ms/step

```
In [ ]: sns.regplot(x=y_test, y=y_pred)
plt.title("Regression Line for Predicted values")
plt.show()
```



```
In [ ]: def regression_metrics_display(y_test, y_pred):  
        print(f"MAE is {metrics.mean_absolute_error(y_test, y_pred)}")  
        print(f"MSE is {metrics.mean_squared_error(y_test, y_pred)}")  
        print(f"R2 score is {metrics.r2_score(y_test, y_pred)}")
```

```
In [ ]: regression_metrics_display(y_test, y_pred)
```

```
MAE is 3.7487237733953136  
MSE is 27.631791791938596  
R2 score is 0.6119240015553091
```